

UNIVERSIDAD TECNOLÓGICA DE BOLÍVAR

PROGRAMA DE INGENIERÍA DE SISTEMAS Y COMPUTACIÓN

PARCIAL 2DO CORTE

Prediction of Susceptibility to Suffer Diabetes

Autor:
Pablo Andres Pertuz Duran T00049795

Professor:
Edwin Alexander Puertas Del Castillo

INTELIGENCIA ARTIFICIAL (ISCO-A06A)

Cartagena, Bolívar
21 de abril de 2022

Parcial 2do Corte | | Inteligencia Artificial ISCO-A06A

April 21, 2022

```
[1]: #Se vincula una cuenta de google drive para importar los dataset

# Mount your google drive in google colab
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

1 2. Data Understanding

Number of Instances: 614 (train), 5 (test)

Number of Attributes: 8 plus class

Columns Description:

1. Number of times pregnant
2. Plasma glucose concentration 2 hours in an oral glucose tolerance test
3. Diastolic blood pressure (mm Hg)
4. Triceps skin fold thickness (mm)
5. 2-Hour serum insulin (mu U/ml)
6. Body mass index (weight in kg/(height in m)²)
7. Diabetes pedigree function
8. Age (years)
9. Class variable (0 or 1)

Class Distribution: (class value 1 is interpreted as “tested positive for diabetes”)

```
[2]: import pandas as pd
```

```
[3]: #import train and test splits

df_train = pd.read_csv("/content/drive/MyDrive/parcial2/train_diabetes.csv")
df_test = pd.read_csv("/content/drive/MyDrive/parcial2/test_diabetes.csv")
```

```
[4]: df_train.head()
```

```
[4]:   p_id  no_times_pregnant  glucose_concentration  blood_pressure  \
0    316                  2                   112                68
```

1	25	11	143	94
2	710	2	93	64
3	658	1	120	80
4	542	3	128	72

	skin_fold_thickness	serum_insulin	bmi	diabetes	pedigree	age	diabetes
0	22	94	34.1		0.315	26	0
1	33	146	36.6		0.254	51	1
2	32	160	38.0		0.674	23	1
3	48	200	38.9		1.162	41	0
4	25	190	32.4		0.549	27	1

```
[5]: #evaluating the equality of the data
df_train.info()

#There are not missing values
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   p_id                  614 non-null   int64
1   no_times_pregnant     614 non-null   int64
2   glucose_concentration 614 non-null   int64
3   blood_pressure        614 non-null   int64
4   skin_fold_thickness   614 non-null   int64
5   serum_insulin         614 non-null   int64
6   bmi                   614 non-null   float64
7   diabetes_pedigree     614 non-null   float64
8   age                   614 non-null   int64
9   diabetes              614 non-null   int64
dtypes: float64(2), int64(8)
memory usage: 48.1 KB
```

```
[6]: #We can also check the uniqueness of each column in the DataFrame

for col in df_train.columns:
    print(col + ": ", len(df_train[col].unique()))
```

```
p_id: 614
no_times_pregnant: 17
glucose_concentration: 127
blood_pressure: 44
skin_fold_thickness: 50
serum_insulin: 163
bmi: 237
diabetes_pedigree: 439
```

age: 52
diabetes: 2

```
[7]: #Check statistics of the data
```

```
df_train.describe()
```

```
[7]:
```

	p_id	no_times_pregnant	glucose_concentration	blood_pressure \
count	614.000000	614.000000	614.000000	614.000000
mean	385.773616	3.853420	120.542345	68.765472
std	223.603024	3.358126	31.252286	19.914836
min	1.000000	0.000000	0.000000	0.000000
25%	191.250000	1.000000	99.000000	62.000000
50%	387.000000	3.000000	117.000000	72.000000
75%	572.750000	6.000000	139.000000	80.000000
max	768.000000	17.000000	197.000000	114.000000

	skin_fold_thickness	serum_insulin	bmi	diabetes pedigree \
count	614.000000	614.000000	614.000000	614.000000
mean	20.244300	79.355049	31.909935	0.466342
std	15.886083	117.709950	8.007699	0.331090
min	0.000000	0.000000	0.000000	0.078000
25%	0.000000	0.000000	27.300000	0.240250
50%	23.000000	17.000000	32.000000	0.361000
75%	32.000000	126.000000	36.600000	0.613500
max	63.000000	846.000000	59.400000	2.420000

	age	diabetes
count	614.000000	614.000000
mean	33.325733	0.348534
std	11.929569	0.476895
min	21.000000	0.000000
25%	24.000000	0.000000
50%	29.000000	0.000000
75%	41.000000	1.000000
max	81.000000	1.000000

```
[8]: df_train.dtypes
```

```
[8]:
```

p_id	int64
no_times_pregnant	int64
glucose_concentration	int64
blood_pressure	int64
skin_fold_thickness	int64
serum_insulin	int64
bmi	float64
diabetes pedigree	float64
age	int64

```
diabetes          int64
dtype: object
```

```
[9]: #Classs distribution

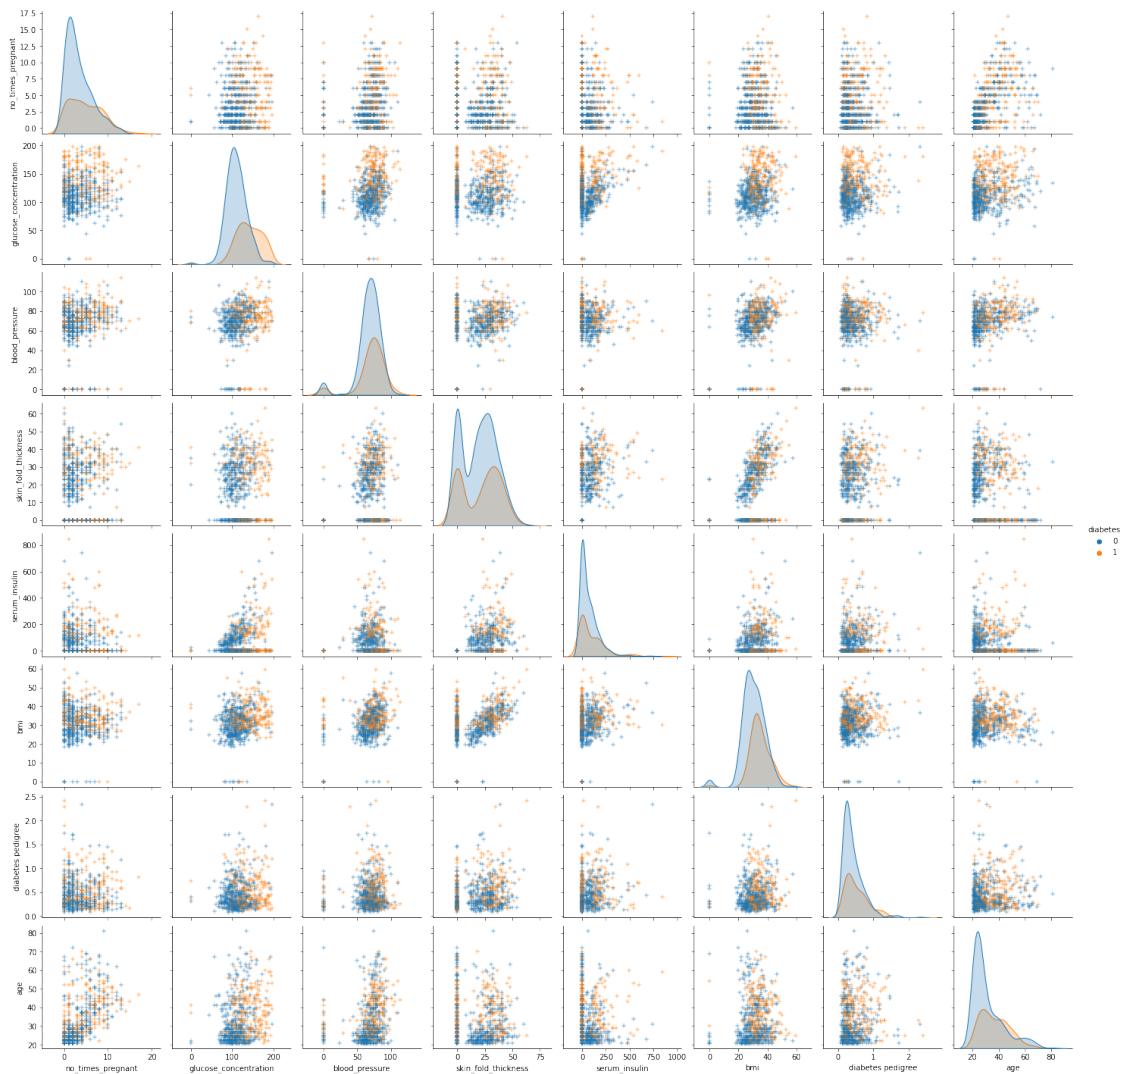
df_train['diabetes'].value_counts()
```

```
[9]: 0    400
     1    214
     Name: diabetes, dtype: int64
```

1.1 Visualize Data

```
[10]: from matplotlib import pyplot as plt
      import seaborn as sns
```

```
[12]: tmp = df_train.drop('p_id', axis=1)
      g = sns.pairplot(tmp, hue='diabetes', markers='+')
      plt.show()
```



Assumption #2 “Y is distributed normally at each value of X” is true for this data

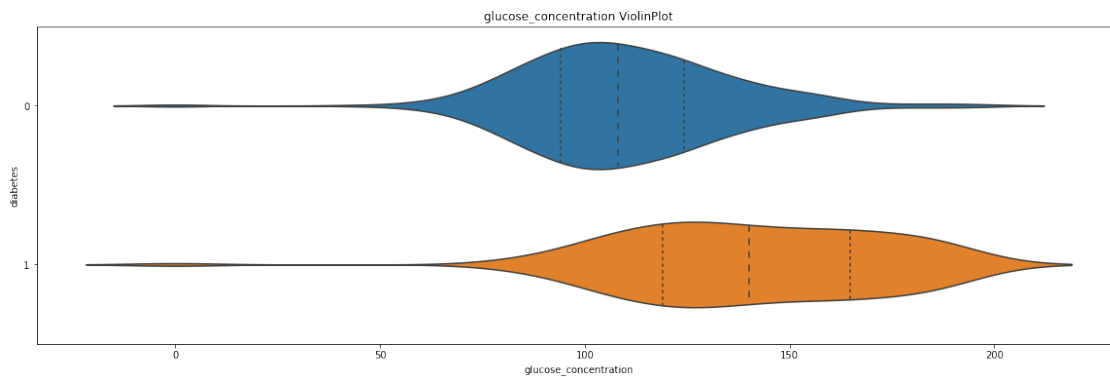
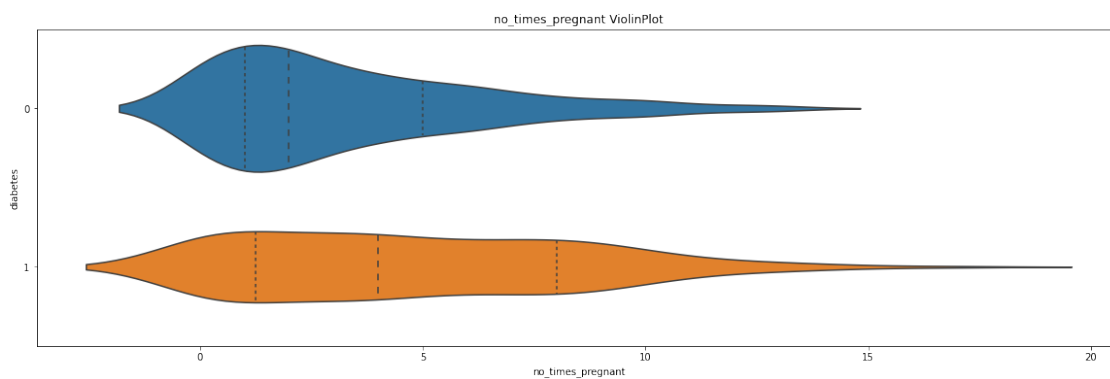
```
[13]: list(df_train.keys())
```

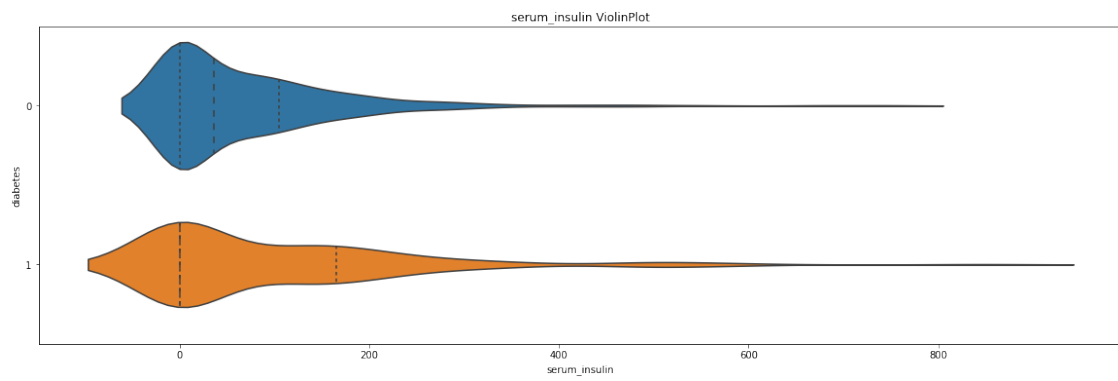
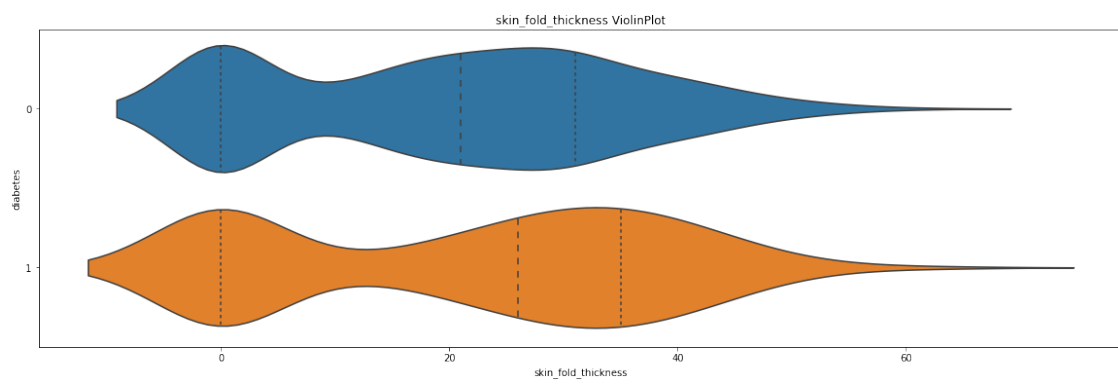
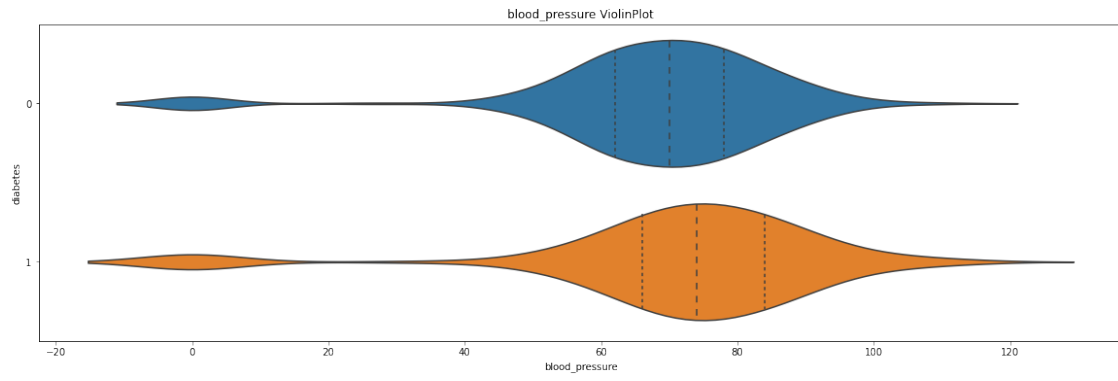
```
[13]: ['p_id',
      'no_times_pregnant',
      'glucose_concentration',
      'blood_pressure',
      'skin_fold_thickness',
      'serum_insulin',
      'bmi',
      'diabetes_pedigree',
      'age',
      'diabetes']
```

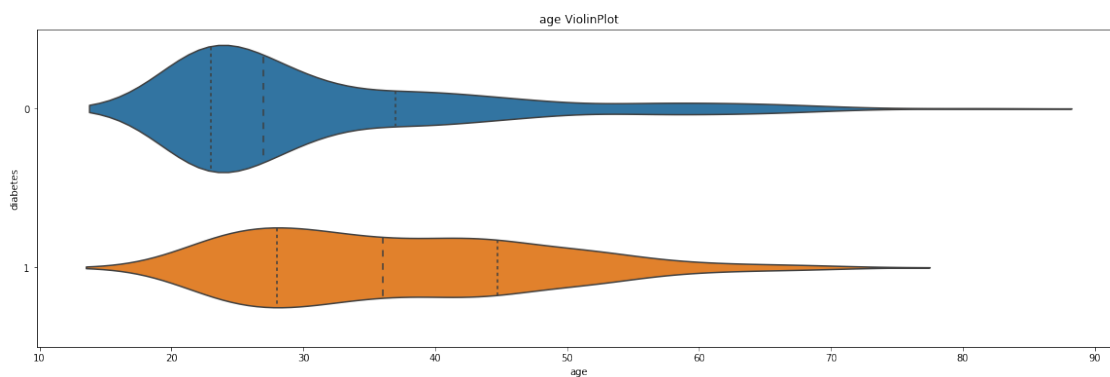
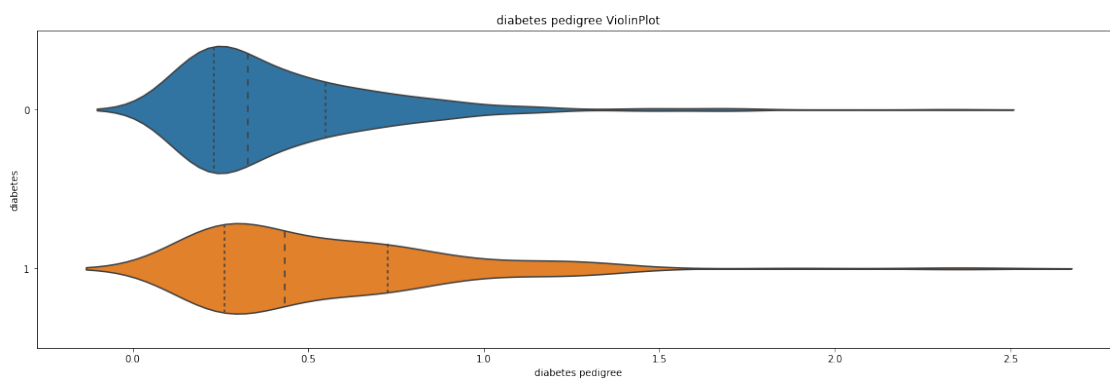
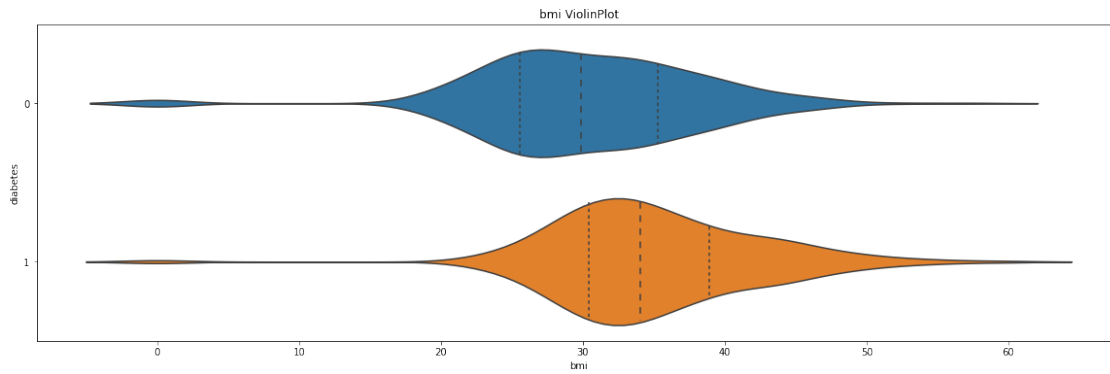
```
[14]: #Converts 0, 1 diabetes classes to str for better visualization
df_temp = df_train.copy()
df_temp['diabetes'] = df_temp['diabetes'].apply(str)
```

```
[15]: #Violin plot
X_col_names = list(df_train.keys())[1:-1]

for col_name in range(len(X_col_names)):
    plt.figure(figsize=(20,6))
    sns.violinplot(x=X_col_names[col_name], y='diabetes', data=df_temp,
        inner='quartile').set_title(X_col_names[col_name]+" ViolinPlot")
    plt.show()
```








It does not seem that any attribute has a strong influence on the determination of diabetes class.

- **Assumption #3** "The variance of Y at every value of X is the same (homogeneity of variances)" seems to be true for this data
- **Assumption #4** "The observations are independent" is true for this data

```
[18]: #Correlation between features
```

```
sns.clustermap(df_train.corr(), cmap = "vlag", dendrogram_ratio = (0.1, 0.2),
    ↳annot = True, linewidths = .8, figsize = (12,13))
plt.show()
```



2 3. Data Preparation

```
[20]: #Check data types
```

```
df_train.dtypes
```

```
[20]: p_id                int64
      no_times_pregnant  int64
      glucose_concentration  int64
      blood_pressure     int64
      skin_fold_thickness int64
      serum_insulin       int64
      bmi                float64
      diabetes pedigree   float64
      age                int64
      diabetes            int64
      dtype: object
```

```
[21]: df_test.dtypes
```

```
[21]: p_id                int64
      no_times_pregnant  int64
      glucose_concentration  int64
      blood_pressure     int64
      skin_fold_thickness int64
      serum_insulin       int64
      bmi                float64
      diabetes pedigree   float64
      age                int64
      dtype: object
```

Both dataframes has correct dtype for values. There are not missing values.

Dataset does not contain categorical features, son one-hot encoding is not necessary.

3 4. Modeling

Note: Cells commented must be ignored.

```
[39]: from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import MinMaxScaler
```

```
[177]: #Define independent and dependent variable:
```

```
X = df_train.drop(['diabetes', 'p_id'], axis=1)
y = df_train['diabetes']
```

3.1 Scaling Features

Scaling numeric features ensures that no particular feature has a disproportionate impact on the model's loss. Optimization algorithms also work better in practice with smaller numbers.

```
[194]: X.describe()
```

```
[194]:
```

	no_times_pregnant	glucose_concentration	blood_pressure	\
count	614.000000	614.000000	614.000000	
mean	3.853420	120.542345	68.765472	
std	3.358126	31.252286	19.914836	
min	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	
50%	3.000000	117.000000	72.000000	
75%	6.000000	139.000000	80.000000	
max	17.000000	197.000000	114.000000	

	skin_fold_thickness	serum_insulin	bmi	diabetes	pedigree	\
count	614.000000	614.000000	614.000000		614.000000	
mean	20.244300	79.355049	31.909935		0.466342	
std	15.886083	117.709950	8.007699		0.331090	
min	0.000000	0.000000	0.000000		0.078000	
25%	0.000000	0.000000	27.300000		0.240250	
50%	23.000000	17.000000	32.000000		0.361000	
75%	32.000000	126.000000	36.600000		0.613500	
max	63.000000	846.000000	59.400000		2.420000	

	age
count	614.000000
mean	33.325733
std	11.929569
min	21.000000
25%	24.000000
50%	29.000000
75%	41.000000
max	81.000000

3.2 Comparing

```
[195]: # Scale to [0, 1]
max_ = X.max(axis=0)
min_ = X.min(axis=0)
X = (X - min_) / (max_ - min_)
```

```
[196]: X.describe()
```

```
[196]:
```

	no_times_pregnant	glucose_concentration	blood_pressure	\
count	614.000000	614.000000	614.000000	

mean	0.226672	0.611890	0.603206
std	0.197537	0.158641	0.174692
min	0.000000	0.000000	0.000000
25%	0.058824	0.502538	0.543860
50%	0.176471	0.593909	0.631579
75%	0.352941	0.705584	0.701754
max	1.000000	1.000000	1.000000

	skin_fold_thickness	serum_insulin	bmi	diabetes	pedigree \
count	614.000000	614.000000	614.000000	614.000000	
mean	0.321338	0.093800	0.537204	0.165816	
std	0.252160	0.139137	0.134810	0.141371	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.459596	0.069278	
50%	0.365079	0.020095	0.538721	0.120837	
75%	0.507937	0.148936	0.616162	0.228651	
max	1.000000	1.000000	1.000000	1.000000	

	age
count	614.000000
mean	0.205429
std	0.198826
min	0.000000
25%	0.050000
50%	0.133333
75%	0.333333
max	1.000000

3.3 Import Dependencies

```
[214]: from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression, LinearRegression
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier

from sklearn.model_selection import cross_val_score

from sklearn.metrics import mean_squared_error, mean_squared_log_error, \
    →accuracy_score
```

```
[256]: linreg_pipeline = Pipeline(steps=[('model', LinearRegression())])
```

```
[257]: logreg_pipeline = Pipeline(steps=[('model', \
    →LogisticRegression(solver='liblinear'))])
```

```
[278]: rf_pipeline = Pipeline(steps=[('model', RandomForestClassifier(n_estimators=200, \
    →random_state=64))])
```

4 5. Evaluating Models (k-Fold Cross Validation)

4.1 Mean Squared Error (MSE)

Linear Regression

```
[270]: # Multiply by -1 since sklearn calculates *negative* MSE
scores = -1 * cross_val_score(linreg_pipeline, X, y,
                              cv=10,
                              scoring='neg_mean_squared_error')

print("MSE scores:\n", scores)
print("\n Average MSE score (across experiments):")
print(scores.mean())
```

MSE scores:

```
[0.1582769  0.15118501 0.13235929 0.17008284 0.17367299 0.16595435
 0.17022733 0.161198   0.18887061 0.18736717]
```

Average MSE score (across experiments):
0.16591944838299763

Logistic Regression

```
[260]: # Multiply by -1 since sklearn calculates *negative* MSE
scores = -1 * cross_val_score(logreg_pipeline, X, y,
                              cv=10,
                              scoring='neg_mean_squared_log_error')

print("MLSE scores:\n", scores)
print("\n Average MLSE score (across experiments):")
print(scores.mean())
```

MLSE scores:

```
[0.10074015 0.10074015 0.08524166 0.12398787 0.12602046 0.08663907
 0.14177302 0.11814418 0.13389674 0.10239163]
```

Average MLSE score (across experiments):
0.11195749369807266

Random Forest for classification

```
[261]: # Multiply by -1 since sklearn calculates *negative* MSE
scores = -1 * cross_val_score(rf_pipeline, X, y,
                              cv=10,
                              scoring='neg_mean_squared_log_error')

print("MLSE scores:\n", scores)
print("\n Average MLSE score (across experiments):")
print(scores.mean())
```

MLSE scores:

```
[0.11623863 0.12398787 0.07749242 0.10074015 0.10239163 0.09451535
0.10239163 0.09451535 0.12602046 0.11814418]
```

Average MLSE score (across experiments):
0.10564376688904713

The problem addressed is of the classification type. Therefore, a linear regression is not the right approach. To solve classification problems, models such as Logistic Regression or Random Forest for Classification are used. We will now compare the last two with the metrics corresponding to classification problems.

4.2 Accuracy

Logistic Regression

```
[275]: scores = cross_val_score(logreg_pipeline, X, y,
                                cv=10,
                                scoring='accuracy')

print("Accuracy scores:\n", scores)
print("\n Average Accuracy score (across experiments):")
print(scores.mean())
```

Accuracy scores:

```
[0.79032258 0.79032258 0.82258065 0.74193548 0.73770492 0.81967213
0.70491803 0.75409836 0.72131148 0.78688525]
```

Average Accuracy score (across experiments):
0.7669751454257007

Random Forest for classification

```
[279]: scores = cross_val_score(rf_pipeline, X, y,
                                cv=10,
                                scoring='accuracy')

print("Accuracy scores:\n", scores)
print("\n Average Accuracy score (across experiments):")
print(scores.mean())
```

Accuracy scores:

```
[0.74193548 0.74193548 0.83870968 0.80645161 0.7704918 0.80327869
0.78688525 0.83606557 0.73770492 0.78688525]
```

Average Accuracy score (across experiments):
0.7850343733474352

4.3 Logarithmic Loss

Logistic Regression

```
[264]: scores = -1 * cross_val_score(logreg_pipeline, X, y,
                                     cv=10,
                                     scoring='neg_log_loss')

print("Logarithmic Loss scores:\n", scores)
print("\n Average Logarithmic Loss score (across experiments):")
print(scores.mean())
```

Logarithmic Loss scores:

```
[0.53029444 0.49533375 0.48809391 0.52422738 0.55542863 0.44782837
 0.52346336 0.50469808 0.60778733 0.4723959 ]
```

Average Logarithmic Loss score (across experiments):
0.5149551168821109

Random Forest for classification

```
[265]: scores = -1 * cross_val_score(rf_pipeline, X, y,
                                     cv=10,
                                     scoring='neg_log_loss')

print("Logarithmic Loss scores:\n", scores)
print("\n Average Logarithmic Loss score (across experiments):")
print(scores.mean())
```

Logarithmic Loss scores:

```
[0.53152699 0.51647937 0.44047989 0.45198319 0.49193008 0.4644009
 0.46463281 0.40011293 0.52648169 0.4149656 ]
```

Average Logarithmic Loss score (across experiments):
0.47029934461234896

4.4 F1 Score

Logistic Regression

```
[266]: scores = cross_val_score(logreg_pipeline, X, y,
                                 cv=10,
                                 scoring='f1')

print("F1 scores:\n", scores)
print("\n Average F1 score (across experiments):")
print(scores.mean())
```

F1 scores:

```
[0.62857143 0.60606061 0.7027027  0.55555556 0.5          0.73170732]
```



```
0.4          0.57142857 0.51428571 0.60606061]
```

Average F1 score (across experiments):
0.5816372501738356

Random Forest for classification

```
[267]: scores = cross_val_score(rf_pipeline, X, y,
                                cv=10,
                                scoring='f1')

print("F1 scores:\n", scores)
print("\n Average F1 score (across experiments):")
print(scores.mean())
```

F1 scores:

```
[0.63414634 0.52941176 0.73684211 0.66666667 0.66666667 0.72727273
 0.62857143 0.7          0.6          0.59459459]
```

Average F1 score (across experiments):
0.6484172295204538

Random forest outperform logistic regression with default hyperparameters.

5 6. Fitting and Evaluating Models with df_test

```
[305]: from sklearn.metrics import confusion_matrix
```

```
[295]: X_full = df_train.drop(['diabetes', 'p_id'], axis=1)
y_full = df_train['diabetes']

# Scale to [0, 1]
max_ = df_train.drop(['diabetes', 'p_id'], axis=1).max(axis=0)
min_ = df_train.drop(['diabetes', 'p_id'], axis=1).min(axis=0)
X_full = (X_full - min_) / (max_ - min_)

#split data into train and validation sets
X_train, X_valid, y_train, y_valid = train_test_split(X_full, y_full,
    ↪test_size=0.3, random_state=64)
```

```
[296]: print(X_train.shape, X_valid.shape, y_train.shape, y_valid.shape)
```

```
(429, 8) (185, 8) (429,) (185,)
```

```
[301]: X_test = df_test.drop('p_id', axis=1)
X_test
```

```
[301]:
```

	no_times_pregnant	glucose_concentration	blood_pressure	\
0	2	112	68	
1	11	143	94	
2	2	93	64	
3	1	120	80	
4	3	128	72	

	skin_fold_thickness	serum_insulin	bmi	diabetes	pedigree	age
0	22	94	34.1		0.315	26
1	33	146	36.6		0.254	51
2	32	160	38.0		0.674	23
3	48	200	38.9		1.162	41
4	25	190	32.4		0.549	27

```
[302]: # Scale to [0, 1]
max_ = df_train.drop(['diabetes', 'p_id'], axis=1).max(axis=0)
min_ = df_train.drop(['diabetes', 'p_id'], axis=1).min(axis=0)
X_test = (X_test - min_) / (max_ - min_)
```

```
[312]: #Function to compute accuracy and confusion matrix for a given inputs and target

def predict_and_plot(inputs, targets, model, name=''):
    preds = model.predict(inputs)

    accuracy = accuracy_score(targets, preds)
    print("Accuracy: {:.2f}%".format(accuracy * 100))

    cf = confusion_matrix(targets, preds, normalize='true')
    plt.figure()
    sns.heatmap(cf, annot=True)
    plt.xlabel('Prediction')
    plt.ylabel('Target')
    plt.title('{} Confusion Matrix'.format(name));
```


5.1 Logistic Regression

```
[324]: #Fitting
logreg_model = LogisticRegression()
logreg_model.fit(X_train, y_train)
```

```
[324]: LogisticRegression()
```


```
[325]: #Predictions and Evaluation (Training set)
predict_and_plot(X_train, y_train, logreg_model, 'Training')
```

Accuracy: 78.32%



```
[326]: #Predictions and Evaluation (Evaluation set)
predict_and_plot(X_valid, y_valid, logreg_model, 'Evaluation')
```

Accuracy: 78.38%



```
[327]: #Predictions on df_test  
logreg_model.predict(X_test)
```

```
[327]: array([0, 1, 0, 0, 0])
```


5.2 RandomForest Classifier

```
[373]: #Fitting  
rf_model = RandomForestClassifier(n_estimators=10, random_state=64)  
rf_model.fit(X_train, y_train)
```

```
[373]: RandomForestClassifier(n_estimators=10, random_state=64)
```

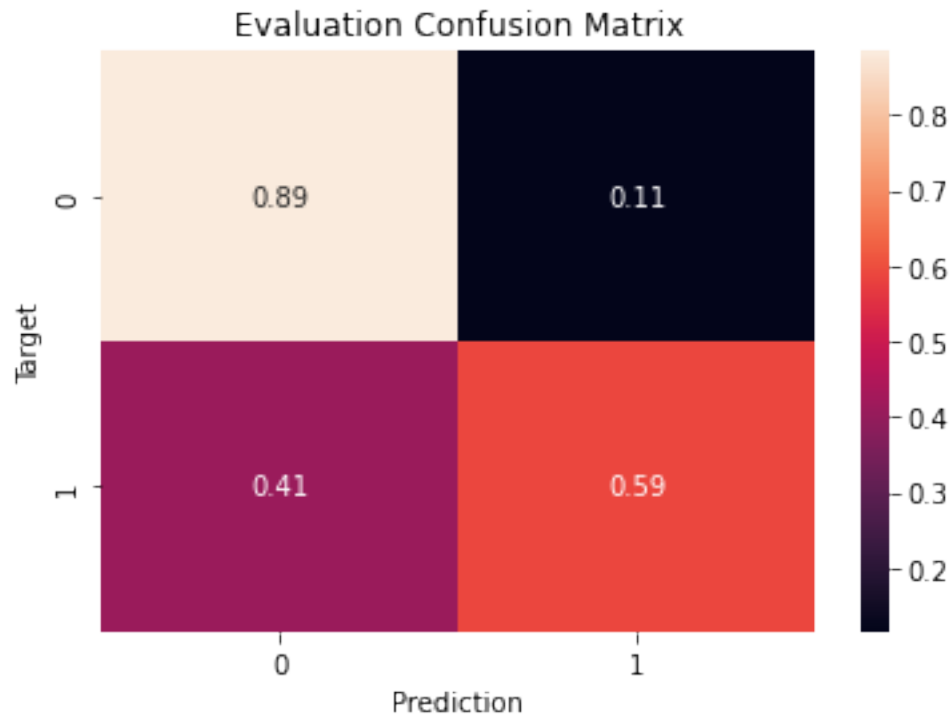
```
[374]: #Predictions and Evaluation (Training set)  
predict_and_plot(X_train, y_train, rf_model, 'Training')
```

Accuracy: 99.07%



```
[375]: #Predictions and Evaluation (Evaluation set)  
predict_and_plot(X_valid, y_valid, rf_model, 'Evaluation')
```

Accuracy: 78.38%



```
[376]: #Predictions on df_test  
rf_model.predict(X_test)
```

```
[376]: array([0, 1, 0, 0, 1])
```

6 7. Conclusion

None of the trained models exceeded an accuracy of 80%. This may be due to the small size of the dataset. It is also necessary to have knowledge of the hyperparameters of each model to choose the ones that improve the performance for the problem.